

Machine Learning Progress Report

Paula Lozano Gonzalo

April 14, 2025

1 Model Overview

This report documents my progress in developing a classifier for diabetes prediction. After identifying significant class imbalance issues in my previous work, I have now implemented a balanced dataset approach and evaluated three model types:

- SGDClassifier (baseline)
- Random Forest classifier (current best performer)
- RidgeClassifierCV (new linear model for comparison)

The task remains binary classification - predicting diabetes outcomes based on patient characteristics.

2 Dataset & Preprocessing

Key updates to the dataset and preprocessing:

- **Class Balancing:** I made a new file to make sure I cleaned the data for equality in between negative and positive cases. And this is how the data is looking now.

```
./clean_data.py
Reading data from data/diabetes_dataset.csv...
Original dataset:
Positive cases: 8500
Negative cases: 91500

Balanced dataset:
Total cases: 17000
Positive cases: 8500
Negative cases: 8500
```

As we can see, I now have 17,000 cases versus the original 100,000. However, I don't think that this is a problem since we still have plenty of data for the model to learn trends.

- **Data Splitting:** I now have to split the data again and here is how it's looking.

```
Original data shape: (17000, 17)
Test data shape: (3400, 17)
Training data shape: (10880, 17)
Validation data shape: (2720, 17)
```

- Maintained previous preprocessing:
 - Missing value imputation (most frequent for categorical, median for numerical)
 - Binary encoding for categorical features
 - Polynomial features (degree=2) where applicable

3 Training Progress

After balancing the dataset, I wanted to re-train the two past models to see if there would be a difference and all models showed improved ability to identify positive cases, though with slightly reduced overall accuracy:

Model	Training Acc.	Validation Acc.	Validation F1
SGDClassifier	0.881	0.890	0.827
Random Forest	0.899	0.898	0.907
RidgeClassifierCV	0.881	0.887	0.879

Table 1: Model performance on balanced dataset

Confusion Matrices (Validation Data):

==== SGDClassifier ====	==== Random Forest ====	==== RidgeClassifierCV ====
t/p F T	t/p F T	t/p F T
F 1140.0 242.0	F 1258.0 124.0	F 1231.0 151.0
T 223.0 1115.0	T 126.0 1212.0	T 171.0 1167.0
Precision: 0.822	Precision: 0.907	Precision: 0.885
Recall: 0.833	Recall: 0.906	Recall: 0.872
F1: 0.827	F1: 0.907	F1: 0.879

The best model so far is the Random Forest so I will be talking about this one mostly to make the report as short as I can but also showing all the work I'm doing.




Figure 1: Precision-Recall Plot for Random Forest (Balanced Data)




Figure 2: Precision-Recall Curve for Random Forest (Balanced Data)

Key observations:

- Random Forest emerges as the best performer ($F1=0.907$) on balanced data
- All models show much better balance between precision and recall compared to imbalanced data (As seen in the confusion matrices)

- The number of false negatives decreased significantly:
 - Random Forest: 126 (vs 240 with imbalanced data)
 - SGDCClassifier: 223 (vs 1990 with imbalanced data)
- Training and validation scores are now much closer, suggesting reduced overfitting

By all these improvements I can tell that the decision of cleaning the data was the correct one and I will stick to it.

4 Challenges & Solutions

Challenge 1: Reduced Dataset Size

- Solution: Accepted smaller dataset (17,000 vs original 100,000) to achieve class balance
- Impact: Models perform well despite smaller size, suggesting quality over quantity

Challenge 2: Model Selection

- Solution: Compared three model types
- Outcome: Random Forest performs best, but linear models remain competitive

Challenge 3: Metric Interpretation

- Solution: Focused on F1 score rather than accuracy as primary metric
- Benefit: Better captures model performance on both classes

5 Next Steps

- Keep messing around with the hyper-parameters to see if I can improve random forest.
- Train neural networks as well and compare with the past models.